



# LaVR: Scene Latent Conditioned Generative Video Trajectory Re-Rendering using Large 4D Reconstruction Models

Mingyang Xie Numair Khan Tianfu Wang Naina Dhirgra Seonghyeon Nam Haitao Yang Zhuo Hui  
Christopher Metzler Andrea Vedaldi Hamed Pirsiavash Lei Luo

Generate views of a **dynamic scene** from a novel camera trajectory




Mingyang Xie [WeChat ID: mingyangxie](#)

Let's Chat!

5th year PhD Student @ UMD  
Email: mingyang@umd.edu

Seeking Full time roles in 2027.





Tianfu Wang

[WeChat ID: Deefyfe](#)

2nd year PhD Student @ UMD  
Webpage: <https://tianfwang.github.io/>  
Looking for Summer Internship in 2027.


About LaVR,  
Research, or  
Anything!



## Our Motivations


- Prior works guide video generation with re-rendered point clouds, which imposes an explicit and relatively hard geometric condition.
- However, even small depth-level errors in point clouds can appear as severe geometric artifacts under novel viewpoints, which then propagate into video generation.
- We instead use 4D scene latents from a Large 4D Reconstruction Model as a softer and more abstract geometric condition for video generation.
- This lets the pretrained video diffusion prior regularize imperfect geometry/depths into more realistic outputs.

## What is a Large 4D Reconstruction Model?




Such a model predicts the latent representation of a 4D scene at every timestep, from which it further predicts the point maps and camera poses.

\*Model architecture of CUT3R. Figure by Wang Qianqian et al. CVPR 2025. More at their webpage <https://cut3r.github.io/>.



Compared to ours, point cloud based methods are prone to depth estimation errors, methods with no 3D conditioning are prone to shape deformation.



Details of Our 4D Scene Latent Conditioning.